

# Catapult Matchlib Memory Logs and Debug Methodology

Stuart Swan

Platform Architect

Siemens EDA

31 January 2023

## Introduction

This document describes the Matchlib memory logging feature and related memory debug methodology. Memories are commonly used in HLS designs, and design and verification engineers often need visibility into what is occurring in the memories in order to be able to debug issues both in pre-HLS scenarios, as well as debugging memory-related problems that may occur when comparing pre-HLS versus post-HLS results.

The memory logging and debug methodology described here is complimentary to the Matchlib channel logs and debug methodology. The channel logs are described in detail in the document titled `matchlib_soc_debug_tutorial.pdf` which is in the Catapult Matchlib examples toolkit in the "doc" directory. It is recommended that you read that document in addition to this document to understand the overall logging and debug methodology.

The memory logs and debug methodology described in this document are not tightly coupled to either Matchlib or SystemC designs, and the same approach can be used in Catapult C++ models.

## Background

Memories appear as normal C/C++ arrays (or similar classes such as `std::array<>`) within HLS models, and are mapped to RAMs during HLS. Pointers may appear in HLS models, and pointer accesses are resolved to array accesses during HLS. Array accesses are coded in models using the indexing operator, for example:

```
uint16 my_mem[128];  
my_mem[x] = y;  
x = my_mem[y];
```

There are several aspects of arrays which complicate debugging when using HLS:

1. HLS often adds latency when it synthesizes the design, so an array access which occurs at a particular simulation time in the pre-HLS model most likely will not occur at the same time in the post-HLS model.
2. Array accesses may be re-ordered in the pre-HLS simulation, and in some cases that is OK, and in other cases it can cause a bug. For example, when a single pre-HLS model operates with and without random stall injection, two processes which are writing to a shared RAM may cause the writes to the RAM to be seen in different orders in the two scenarios. If the two processes always write to independent addresses, there never is a bug. However, if the two processes

write to the same addresses, then a "WRITE/WRITE" conflict is possible and different overall simulation results may be seen in the two scenarios.

3. Array accesses may be re-ordered by HLS. For example, HLS may reorder READ operations (which is typically safe to do), and may reorder WRITE operations in certain situations. Certain HLS directives such as "ignore\_memory\_precedences" explicitly allow HLS to do reordering of accesses.

The memory logs and debug methodology described in this document are designed to address the above issues.

## Memory Logs

The memory log feature is provided by a few classes that add the logging functionality. The most commonly used class is `extended_array`:

```
#include <array>

template <typename T, unsigned N>
class extended_array : public std::array<T,N> { .. }
```

This class extends `std::array<>`, which itself is a simple wrapper around a normal C array, i.e.

```
T my_array[N];
```

The `extended_array` template class overloads the indexing operators "`[]`" to gain visibility into when array reads and writes are performed in the pre-HLS model. It then generates log files to record the array accesses. For each array instance, a separate log file is generated for all the read operations, and another log is generated for the write operations. All log files are in ASCII, and all numbers are represented in hex format, and each read/write operation takes exactly one line. The format of the log files is:

```
Address  write_count  value
```

For example, in a "read" log file, the following entry:

```
0000040e 2 000000000000000070
```

indicates that address 0x40e was read and the read value was 0x70, and that read value represents what was written on the second write operation to that location in the RAM.

Note that both "read" logs and "write" logs have write counts in them, there is no corresponding "read count".

As we will see, the "write count" is a key tool to enable effective debugging strategies even in the presence of the complicating factors mentioned above.

The constructor for `extended_array` is:

```
extended_array(std::string _nm = "", bool _add_time_stamp=0) { ...}
```

If the first argument is omitted or an empty string, then no log will be produced, else the argument is used to name the read and write log files for that array instance. If the second argument is true, then the simulation time will be appended to each line in the log files. This can be helpful in certain debugging situations, however by default this option is not enabled because it causes "diff" and related comparison utilities to flag harmless differences.

The `extended_array` class is not the only way to generate memory logs. In example 80\* in the Matchlib toolkit, as we will see below, we can also generate these logs with the same format from RTL models. It is also possible to add this logging format to customer-specific memory models. Once that is done, the methodology described in this document can be easily used.

### Example #1

Let's first see how memory logs are used in a simple example, which is example 12\* in the Matchlib examples toolkit. This example is a simple "ping-pong" memory that is written by a producer process and read by a consumer process. The two processes share a memory that has 16 entries. The producer writes to the first set of eight entries in ascending address order, then hands ownership of that half over to the consumer process, which reads the entries in ascending address order. Once the producer has handed ownership of a half of memory over, it proceeds to write to the other half of the memory. The ping-ping handoff sequence then repeats.

You can build and run the example by doing:

```
cd 12_ping_pong_mem
make build
./sim_sc
```

This will produce the files:

```
mem_prehls_write.log
mem_prehls_read.log
```

In this case the "write.log" is what the producer process has written, and the "read.log" is what the consumer process has read. If we use the "kompare" utility to compare the logs, we will see:

mem_prehls_write.log	mem_prehls_read.log
1 00000000 1 00000	1 00000000 1 00000
2 00000001 1 00001	2 00000001 1 00001
3 00000002 1 00002	3 00000002 1 00002
4 00000003 1 00003	4 00000003 1 00003
5 00000004 1 00004	5 00000004 1 00004
6 00000005 1 00005	6 00000005 1 00005
7 00000006 1 00006	7 00000006 1 00006
8 00000007 1 00007	8 00000007 1 00007
9 00000008 1 00008	9 00000008 1 00008
10 00000009 1 00009	10 00000009 1 00009
11 0000000a 1 0000a	11 0000000a 1 0000a
12 0000000b 1 0000b	12 0000000b 1 0000b
13 0000000c 1 0000c	13 0000000c 1 0000c
14 0000000d 1 0000d	14 0000000d 1 0000d
15 0000000e 1 0000e	15 0000000e 1 0000e
16 0000000f 1 0000f	16 0000000f 1 0000f
17 00000000 2 00010	17 00000000 2 00010
18 00000001 2 00011	18 00000001 2 00011
19 00000002 2 00012	19 00000002 2 00012
20 00000003 2 00013	20 00000003 2 00013
21 00000004 2 00014	21 00000004 2 00014
22 00000005 2 00015	22 00000005 2 00015
23 00000006 2 00016	23 00000006 2 00016
24 00000007 2 00017	24 00000007 2 00017
25 00000008 2 00018	25 00000008 2 00018
26 00000009 2 00019	26 00000009 2 00019
27 0000000a 2 0001a	27 0000000a 2 0001a
28 0000000b 2 0001b	28 0000000b 2 0001b
29 0000000c 2 0001c	29 0000000c 2 0001c
30 0000000d 2 0001d	30 0000000d 2 0001d
31 0000000e 2 0001e	31 0000000e 2 0001e
32 0000000f 2 0001f	32 0000000f 2 0001f
33 00000000 3 00020	
34 00000001 3 00021	
35 00000002 3 00022	
36 00000003 3 00023	
37 00000004 3 00024	
38 00000005 3 00025	
39 00000006 3 00026	

Comparing file file:///mem\_prehls\_read.log 1 of 1 difference, 0 applied 1 of 1 file

First, we can note the extra entries at the bottom of the write log. These represent write operations to the memory that were never read by the consumer process. The reason the consumer never read them is because the testbench ended the test before the consumer process had sufficient time to read these items. In this case this result is intended by the author of the testbench.

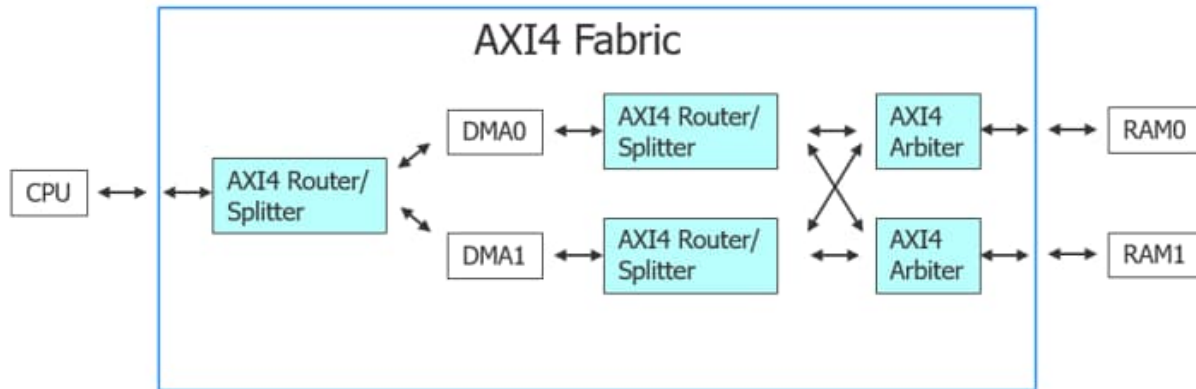
The second item to notice is the write\_count fields. Note that when the producer writes to the same addresses for the second and third times, the write\_count field clearly indicates that this occurred. Note that the identical write\_count values appear in the read log.

Lastly, you should note that in this simple test case the order of write addresses and read addresses are identical. That may not always be the case for other designs, and if the order is different it may not necessarily represent a bug, as we will see.

## Example #2

The next example using memory logs is example 60\* in the Catapult Matchlib toolkit. This example is extensively described in the document titled matchlib\_soc\_debug\_tutorial.pdf which is in the Catapult Matchlib examples toolkit in the "doc" directory. We will be looking at the same tests described in that document, so we will only briefly describe the test here.

The design is an AXI4 fabric consisting of two DMAs and two RAMs. The memory logs are generated from each of the RAMs:



I/\*\*

For this test scenario the DUT performs a single specific test. As per AXI convention, all addresses are byte addresses. The test is as follows: At startup every 64-bit word in the RAMs is initialized with its byte address within that specific RAM. For example, the second word in each RAM will be initialized with the value eight. The CPU will program DMA0 to perform a copy operation that reads from RAM0 address 0x0000 and writes to RAM0 address 0x2000. The copy operation is comprised of 16 AXI4 bursts, where the burst length has been limited to 2 beats per burst. After the above copy operation, the CPU will program DMA1 to perform a similar copy operation from RAM0 address 0x2000 to RAM1 address 0x2000. At the completion of the test we should see that 16 words at RAM0 address 0x0000 have been copied to RAM0 address 0x2000 (by DMA0), and we should also see those same 16 words copied to RAM1 address 0x2000 (by DMA1). The testbench is self-checking and at the completion of the simulation it flags an error if the memory contents are not correct.

You can build the design and run the "no\_stall" and "stall" simulations with the following commands:

```
cd 60_rand_stall
make all
```

This will produce the following memory log files:

```
top.ram0_no_stall_read.log
top.ram0_no_stall_write.log
top.ram0_stall_read.log
top.ram0_stall_write.log
top.ram1_no_stall_read.log
top.ram1_no_stall_write.log
top.ram1_stall_read.log
top.ram1_stall_write.log
```

The "no\_stall" files represent the desired system behavior, and the "stall" files represent the incorrect system behavior that is described in detail in the [matchlib\\_soc\\_debug\\_tutorial.pdf](#) document. Briefly, the incorrect behavior arises in the "stall" scenario because the second copy operation starts too early, so it reads incorrect data from RAM0.

All the addresses in the log files are indexes into the RAM arrays, which contain 8 bytes per element. Since AXI4 addresses are always byte addresses, we need to keep in mind the mapping between AXI4

byte addresses and the array indexes. For example, the byte address 0x2000 needs to be divided by 8 to arrive at the array index address of 0x400.

The result of the incorrect behavior described in the matchlib\_soc\_debug\_tutorial.pdf document is easily seen when comparing the read and write operations on RAM1:

```
kompare *ram1*read*
```

top.ram1_no_stall_read.log	top.ram1_stall_read.log
1 00000400 2 0000000000000000	1 00000400 2 0000000000002000
2 00000401 2 0000000000000008	2 00000401 2 0000000000002008
3 00000402 2 0000000000000010	3 00000402 2 0000000000002010
4 00000403 2 0000000000000018	4 00000403 2 0000000000002018
5 00000404 2 0000000000000020	5 00000404 2 0000000000002020
6 00000405 2 0000000000000028	6 00000405 2 0000000000002028
7 00000406 2 0000000000000030	7 00000406 2 0000000000002030
8 00000407 2 0000000000000038	8 00000407 2 0000000000002038
9 00000408 2 0000000000000040	9 00000408 2 0000000000002040
10 00000409 2 0000000000000048	10 00000409 2 0000000000002048
11 0000040a 2 0000000000000050	11 0000040a 2 0000000000002050
12 0000040b 2 0000000000000058	12 0000040b 2 0000000000002058
13 0000040c 2 0000000000000060	13 0000040c 2 0000000000002060
14 0000040d 2 0000000000000068	14 0000040d 2 0000000000002068
15 0000040e 2 0000000000000070	15 0000040e 2 0000000000002070
16 0000040f 2 0000000000000078	16 0000040f 2 0000000000002078

```
kompare *ram1*write*
```

5537 00000400 2 0000000000000000	5537 00000400 2 0000000000002000
5538 00000401 2 0000000000000008	5538 00000401 2 0000000000002008
5539 00000402 2 0000000000000010	5539 00000402 2 0000000000002010
5540 00000403 2 0000000000000018	5540 00000403 2 0000000000002018
5541 00000404 2 0000000000000020	5541 00000404 2 0000000000002020
5542 00000405 2 0000000000000028	5542 00000405 2 0000000000002028
5543 00000406 2 0000000000000030	5543 00000406 2 0000000000002030
5544 00000407 2 0000000000000038	5544 00000407 2 0000000000002038
5545 00000408 2 0000000000000040	5545 00000408 2 0000000000002040
5546 00000409 2 0000000000000048	5546 00000409 2 0000000000002048
5547 0000040a 2 0000000000000050	5547 0000040a 2 0000000000002050
5548 0000040b 2 0000000000000058	5548 0000040b 2 0000000000002058
5549 0000040c 2 0000000000000060	5549 0000040c 2 0000000000002060
5550 0000040d 2 0000000000000068	5550 0000040d 2 0000000000002068
5551 0000040e 2 0000000000000070	5551 0000040e 2 0000000000002070
5552 0000040f 2 0000000000000078	5552 0000040f 2 0000000000002078

When we do "diff \*ram0\*write\*" we see that there are no differences at all in the two scenarios in terms of what was written to RAM0.

However, when we do:

```
kompare *ram0*read*
```

We see:

top.ram0_no_stall_read.log		top.ram0_stall_read.log
1 00000000 1 000000000000000000		1 00000000 1 000000000000000000
2 00000001 1 000000000000000008		2 00000001 1 000000000000000008
3 00000002 1 000000000000000010		3 00000002 1 000000000000000010
4 00000003 1 000000000000000018		4 00000003 1 000000000000000018
5 00000004 1 000000000000000020		5 00000400 1 000000000000000000
6 00000005 1 000000000000000028		6 00000401 1 000000000000000008
7 00000400 2 000000000000000000		7 00000004 1 000000000000000020
8 00000401 2 000000000000000008		8 00000005 1 000000000000000028
9 00000006 1 000000000000000030		9 00000402 1 000000000000000010
10 00000007 1 000000000000000038		10 00000403 1 000000000000000018
11 00000402 2 000000000000000010		11 00000006 1 000000000000000030
12 00000403 2 000000000000000018		12 00000007 1 000000000000000038
13 00000008 1 000000000000000040		13 00000404 1 000000000000000020
14 00000009 1 000000000000000048		14 00000405 1 000000000000000028
15 00000404 2 000000000000000020		15 00000008 1 000000000000000040
16 00000405 2 000000000000000028		16 00000009 1 000000000000000048
17 0000000a 1 000000000000000050		17 0000000a 1 000000000000000050
18 0000000b 1 000000000000000058		18 0000000b 1 000000000000000058
19 00000406 2 000000000000000030		19 00000406 1 000000000000000030
20 00000407 2 000000000000000038		20 00000407 1 000000000000000038
21 0000000c 1 000000000000000060		21 0000000c 1 000000000000000060
22 0000000d 1 000000000000000068		22 0000000d 1 000000000000000068
23 00000408 2 000000000000000040		23 00000408 1 000000000000000040
24 00000409 2 000000000000000048		24 00000409 1 000000000000000048
25 0000000e 1 000000000000000070		25 0000040a 1 000000000000000050
26 0000000f 1 000000000000000078		26 0000040b 1 000000000000000058
27 0000040a 2 000000000000000050		27 0000040c 1 000000000000000060
28 0000040b 2 000000000000000058		28 0000040d 1 000000000000000068
29 0000040c 2 000000000000000060		29 0000040e 1 000000000000000070
30 0000040d 2 000000000000000068		30 0000040f 1 000000000000000078
31 0000040e 2 000000000000000070		31 0000000e 1 000000000000000070
32 0000040f 2 000000000000000078		32 0000000f 1 000000000000000078
33 00000000 1 000000000000000000		33 00000000 1 000000000000000000
34 00000400 2 000000000000000000		34 00000400 2 000000000000000000
35 00000001 1 000000000000000008		35 00000001 1 000000000000000008
36 00000401 2 000000000000000008		36 00000401 2 000000000000000008
37 00000002 1 000000000000000010		37 00000002 1 000000000000000010
38 00000402 2 000000000000000010		38 00000402 2 000000000000000010

Comparing file:///home/sswan/matchlib\_examples\_kit\_...\_examples/examples/60\_rand\_stall/top.ram0\_stall\_read.log 1 of 7 differences. 0 applied 1 of 1 file

The left side is the correct behavior, and the right side is the "stall" incorrect behavior.

All the reads starting at indexes 0x400 and higher represent the second copy operation in the test scenario. We see on the left that those read operations all read data with "write\_counts" that are consistently 2. However, on the right side, we see the incorrect behavior where those reads are occurring where some write\_counts are still 1. This is the root cause of the bug - it is a clear indication that the reads from the second copy operation are occurring too soon in the stall scenario.

### Example #3

The third example is example 80\* in the Matchlib toolkit. In this example we generate memory logs for the pre-HLS and post-HLS model and compare the results. This is the code in the DUT:

```
typedef ac_int<32, false> elem_type;
static const int mem_size = 128;

#ifdef USE_EXT_ARRAY
    extended_array<elem_type, mem_size> mem{"mem_prehls"};
#else
    elem_type mem[mem_size];
#endif

Connections::In <elem_type> CCS_INIT_S1(raddr_in);
Connections::In <elem_type> CCS_INIT_S1(waddr_in);
Connections::In <elem_type> CCS_INIT_S1(data_in);
Connections::Out<elem_type> CCS_INIT_S1(data_out);
```



```

void main() {
    // ...
    wait(); // Reset state

#pragma hls_pipeline_init_interval 1
#pragma pipeline_stall_mode flush
    while (1) {
        mem[waddr_in.Pop()] = data_in.Pop();
        elem_type t = mem[raddr_in.Pop()];
        data_out.Push(t);
    }
}

```

In this design we are mapping "mem" to a RAM with 1 read and 1 write port, and our goal is to pipeline with an  $II=1$ . In the pre-HLS model source code, the mem write operation occurs strictly before the memory read operation, and the required behavior is clear even if waddr and raddr are the same. In this specific case the required behavior would be that whatever is read from data\_in would be written to data\_out. Catapult by default will want to preserve this strict ordering, and this will prevent it from pipelining this design with an  $II=1$  unless we explicitly tell it to ignore the scheduling precedence for the memory operations. We do this by using the following Catapult directive:

```
ignore_memory_precedences -from *write_mem* -to *read_mem*
```

This then allows Catapult to schedule this design with an  $II=1$ , but the user is effectively pledging to never allow waddr and raddr to be the same in the same iteration, or else the pre-HLS behavior and the post-HLS behavior will mismatch (since the ordering relationship in the pre-HLS model is no longer preserved in the post-HLS model).

You can run the pre-HLS simulation by doing:

```

cd 80_mem_log_pre_post_hls
make build
./sim_sc

```

You can run HLS by doing:

```
catapult -f go_hls.tcl &
```

Launch Questa by typing:

```
dofile scverify.tcl
```

In Questa type:

```

run -all
wave zoom full

```

Look at the waveforms and then properly quit from the simulator. Properly quitting will cause the post-HLS memory logs to be created.



At this point the following memory logs will exist:

```
mem_posthls_read.log
mem_posthls_write.log
mem_prehls_read.log
mem_prehls_write.log
```

There is a memory initialization loop in the reset state of the DUT:

```
void main() {
    for (unsigned u=0; u < mem_size; u++)
        mem[u] = 0;
    wait(); // Reset state
```

In the pre-HLS model, the memory is initialized by writing locations in ascending order, as per the code. Interestingly, in the post-HLS model, the memory is initialized in the reverse order. This behavior is not incorrect, it is just different. However, it does complicate the comparison of memory logs a bit. We need to sort the logs by address to see the real differences:

```
sort -k1 mem_posthls_write.log > mem_posthls_write_sorted.log
sort -k1 mem_prehls_write.log > mem_prehls_write_sorted.log
diff mem_prehls_write_sorted.log mem_posthls_write_sorted.log
```

This will show there are no real differences in what is written to the RAM in the pre-HLS versus post-HLS models.

Now you can compare the read logs:

```
kompere *pre*read* *post*read*
```

This will show:

mem_prehls_read.log	mem_posthls_read.log
1 00000064 1 000000000	1 00000064 1 000000000
2 00000064 1 000000000	2 00000064 1 000000000
3 00000064 1 000000000	3 00000064 1 000000000
4 00000002 2 000000002	4 00000002 2 000000002
5 00000003 2 000000003	5 00000003 2 000000003
6 00000004 2 000000004	6 00000004 2 000000004
7 00000005 2 000000005	7 00000005 2 000000005
8 00000006 2 000000006	8 00000006 2 000000006
9 00000007 2 000000007	9 00000007 2 000000007
10 00000008 2 000000008	10 00000008 2 000000008
11 0000000a 2 00000000a	11 0000000a 2 00000000a
12 0000000a 2 00000000a	12 0000000b 2 00000000b
13 0000000b 2 00000000b	13 0000000c 2 00000000c
14 0000000c 2 00000000c	14 0000000d 2 00000000d
15 0000000d 2 00000000d	15 0000000e 2 00000000e
16 0000000e 2 00000000e	16 0000000f 2 00000000f
17 0000000f 2 00000000f	17 00000010 2 000000010
18 00000010 2 000000010	18 00000011 2 000000011
19 00000011 2 000000011	19 00000012 2 000000012
20 00000012 2 000000012	20 00000013 2 000000013
21 00000013 2 000000013	21 00000014 2 000000014
22 00000014 2 000000014	22 00000015 2 000000015
23 00000015 2 000000015	23 00000016 2 000000016
24 00000016 2 000000016	24 00000017 2 000000017
25 00000017 2 000000017	25 00000018 2 000000018
26 00000018 2 000000018	26 00000019 2 000000019
27 00000019 2 000000019	27 0000001a 2 00000001a
28 0000001a 2 00000001a	28 0000001b 2 00000001b
29 0000001b 2 00000001b	29 0000001c 2 00000001c
30 0000001c 2 00000001c	30 0000001d 2 00000001d
31 0000001d 2 00000001d	31 0000001e 2 00000001e
32 0000001e 2 00000001e	32 0000001f 2 00000001f
33 0000001f 2 00000001f	33 00000020 2 000000020
34 00000020 2 000000020	34 00000021 2 000000021
35 00000021 2 000000021	35 00000022 2 000000022
36 00000022 2 000000022	36 00000023 2 000000023
37 00000023 2 000000023	37 00000024 2 000000024
38 00000024 2 000000024	38 00000025 2 000000025

In this case we see that the read at address 0xa in the pre-HLS simulation is missing in the post-HLS simulation model. Looking at the testbench.cpp file, we see that this is where we are intentionally violating the "ignore\_memory\_precedence" pledge discussed at the beginning of this section. Specifically, we are intentionally reading and writing to the same address in the same cycle. The post-HLS RAM model simply omits the read operation in this error case, so it shows up in the memory logs as a missing read operation.

To conclude the discussion on this example, it should be clear that the memory log feature is a useful tool for comparing pre-HLS and post-HLS simulation results, and it is also useful for detecting and debugging incorrect usage of directives such as "ignore\_memory\_precedence".

#### Example #4

The fourth example is Matchlib toolkit example 41\*. This example uses the `ScratchpadClass` to build a banked memory in the DUT, and is explained in detail in the `matchlib_memory_modeling_methodology.pdf` document within the Matchlib toolkit examples "doc" directory. The `ScratchpadClass` is instrumented to produce memory logs in the same manner as `extended_array`. To see the memory logs that are generated for this example, type:

```
cd 41_scratchpad_class
make build
./sim_sc
```

This will produce the memory logs:

```
mem_prehls_read.log
mem_prehls_write.log
```

The orders of writes and reads are random, so they should be sorted so they are understandable:

```
sort -k1 *write* > write_sorted.log
sort -k1 *read* > read_sorted.log
kompare w*sorted* r*sorted*
```

write_sorted.log	read_sorted.log
2264 00002fe7 1 00a5a9af4	3035 00002ff2 1 03ced50e7
2265 00002fe8 1 02cbd7e7c	3036 00002ff3 1 0486022f2
2266 00002fe9 1 02769f861	3037 00002ff3 1 0486022f2
2267 00002fea 1 03a1d58d8	3038 00002ff4 1 0544ba095
2268 00002feb 1 03fe400f7	3039 00002ff4 1 0544ba095
2269 00002fec 1 0685f9917	3040 00002ff5 1 05d9dffce
2270 00002fed 1 062ac04c5	3041 00002ff5 1 05d9dffce
2271 00002fee 1 0202ec34	3042 00002ff5 1 05d9dffce
2272 00002fef 1 02e3fd41c	3043 00002ff6 1 017188f26
2273 00002ff0 1 01c75aa57	3044 00002ff6 1 017188f26
2274 00002ff1 1 06c25c3aa	3045 00002ff6 1 017188f26
2275 00002ff2 1 03ced50e7	3046 00002ff7 1 006c772d0
2276 00002ff3 1 0486022f2	3047 00002ff7 1 006c772d0
2277 00002ff4 1 0544ba095	3048 00002ff7 1 006c772d0
2278 00002ff5 1 05d9dffce	3049 00002ff8 1 071008b68
2279 00002ff6 1 017188f26	3050 00002ff8 1 071008b68
2280 00002ff7 1 006c772d0	3051 00002ff8 1 071008b68
2281 00002ff8 1 071008b68	3052 00002ff9 1 01a15729a
2282 00002ff9 1 01a15729a	3053 00002ff9 1 01a15729a
2283 00002ffa 1 00563b857	3054 00002ff9 1 01a15729a
2284 00002ffb 1 04c9e96a8	3055 00002ffa 1 00563b857
2285 00002ffc 1 0721e861e	3056 00002ffa 1 00563b857
2286 00002ffd 1 00ed67cd4	3057 00002ffa 1 00563b857
2287 00002ffe 1 006399430	3058 00002ffb 1 04c9e96a8
2288 00002fff 1 06f3c3b37	3059 00002ffb 1 04c9e96a8
2289 00003000 1 0617c643e	3060 00002ffc 1 0721e861e
2290 00003001 1 076847e27	3061 00002ffc 1 0721e861e
2291 00003002 1 00732fa74	3062 00002ffd 1 00ed67cd4
2292 00003003 1 060cd519a	3063 00002ffd 1 00ed67cd4
2293 00003004 1 020a9ad14	3064 00002ffe 1 006399430
2294 00003005 1 051f85976	3065 00002ffe 1 006399430
2295 00003006 1 06b27ec8e	3066 00002fff 1 06f3c3b37
2296 00003007 1 04d672b90	3067 00002fff 1 06f3c3b37
2297 00003008 1 0796251d8	3068 00003000 1 0617c643e
2298 00003009 1 025454566	3069 00003001 1 076847e27
2299 0000300a 1 00d4b2c87	3070 00003002 1 00732fa74
2300 0000300b 1 0617c643e	3071 00003003 1 060cd519a
2301 0000300c 1 020a9ad14	3072 00003004 1 020a9ad14

Briefly, what this is showing is that all addresses are written exactly once (on the left), and that they are sometimes read multiple times (note duplicated addresses in purple on the right). Some addresses are never read. Even when there are duplicate reads, if you look carefully you can see that the values on the left and right always match. All of these are correct results, reflecting the code in the testbench and DUT.

### General Tips for Using Memory Logs

1. Because all entries in memory write logs and read logs have write counts, log files can be sorted based on their address (`sort -k1`), and key information is not lost. This is valuable since memory access operations may be legitimately reordered in both pre-HLS scenarios (as discussed at the beginning of this document), and memory accesses may also be reordered during HLS.
2. In effect, memory logs treat each address in the memory as an independent communication channel, and the `write_count` indicates the sequence count for the transaction (or "message")

transmitted through that specific channel. The value field in the memory log indicates the value of the transaction written to or read from the channel.

3. The most common access pattern for memory usage is where every value written to a memory address is read exactly once. In this case, the write counts and values will be identical in the write logs versus the read logs.
4. Other patterns are possible, for example every value written to a memory address may always be read twice. This is analogous to an "upsample" operation in the DSP domain, and in terms of the memory logs you would see the read log showing read operations to addresses occurring twice with the same write count.
5. In most HW designs, the memory access patterns are simple and fixed, and should fit well with the framework described above.
6. When using memory logs to debug a problem, usually if the write\_counts differ between the correct log and the incorrect log, this is likely to point to the root cause.