



Scatter Gather DMA Example

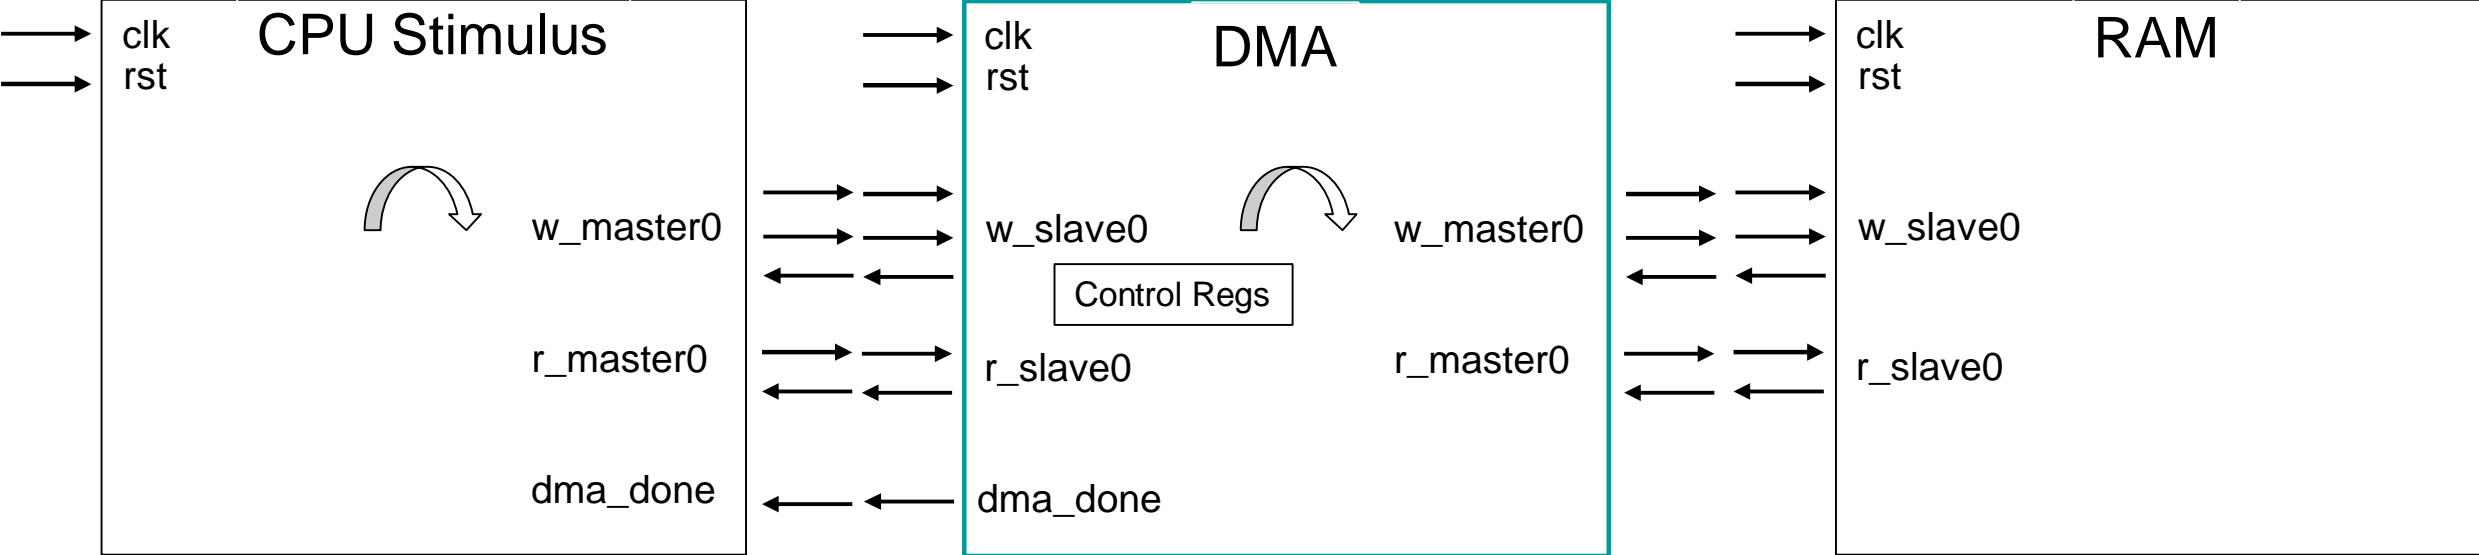
Refer to Matchlib example 18_scatter_gather

Stuart Swan
Platform Architect
Siemens EDA
16 April 2024

Scatter Gather DMA Introduction


- Before reading this presentation, read the “Matchlib AXI4 Interfaces” section within the Matchlib training slides in the Matchlib toolkit doc directory.
- DMA modules are common in SOCs to offload data movement tasks from other blocks
- The Scatter Gather DMA is a simple but representative example of common SOC blocks that use bus fabrics and which are not only datapath blocks but also include control
- Also represents typical DV aspects for SOC blocks
- The DMA fully uses the AXI4 features explained previously such as automatic burst segmentation
- The scatter gather DMA adds features to the simple DMA block presented previously
- Same top level interfaces as simple DMA
- Same RAM block

Scatter Gather DMA block diagram



```
/**
 * * \brief dma module
 */
#pragma hls_design top
class dma : public sc_module, public local_axi {
public:
    sc_in<bool> INIT_S1(clk);
    sc_in<bool> INIT_S1(rst_bar);

    r_master INIT_S1(r_master0);
    w_master INIT_S1(w_master0);
    r_slave INIT_S1(r_slave0);
    w_slave INIT_S1(w_slave0);
    Connections::Out<bool> INIT_S1(dma_done);
```

 = top level of design

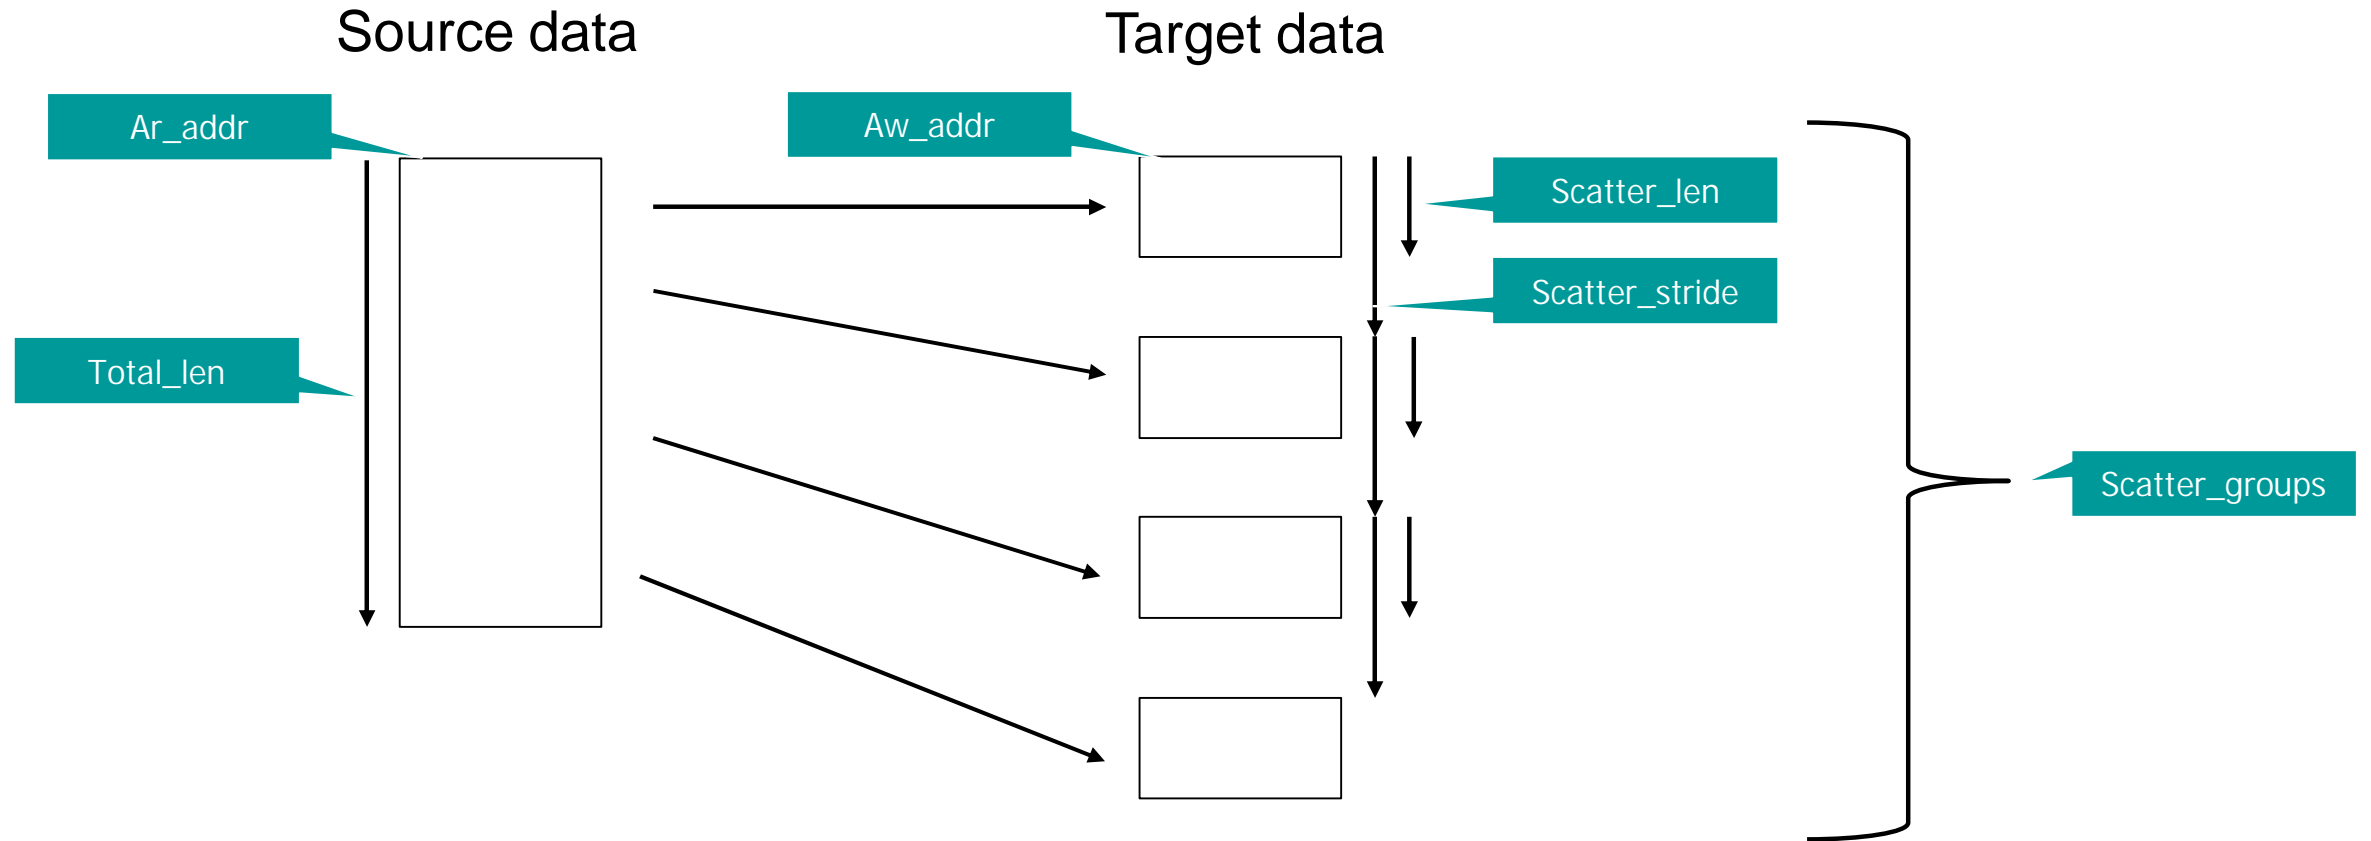
Scatter Gather Slave Address Map

```

6
7 enum dma_mode_t {COPY=0, SCATTER=1, GATHER=2};
8
9 /**
10  *  * \brief dma address map as seen by the CPU
11  */
12 struct dma_address_map
13 {
14     uint64_t  ar_addr;        // source address (byte address as per AXI)
15     uint64_t  aw_addr;        // target address (byte address as per AXI)
16     uint64_t  total_len;      // total length to be copied in bytes
17     uint64_t  scatter_stride; // stride between each scatter group, in bytes
18     uint64_t  scatter_len;    // length of each scatter group, in bytes
19     uint64_t  scatter_groups; // number of scatter groups
20     uint64_t  dma_mode;       // COPY, SCATTER, GATHER
21     uint64_t  start;          // DMA command is complete, cause it to be queued to start
22 };
--
```

Note: In previous DMA (example 08*), total_len was in beats, but now it is a byte length

Scatter Operation

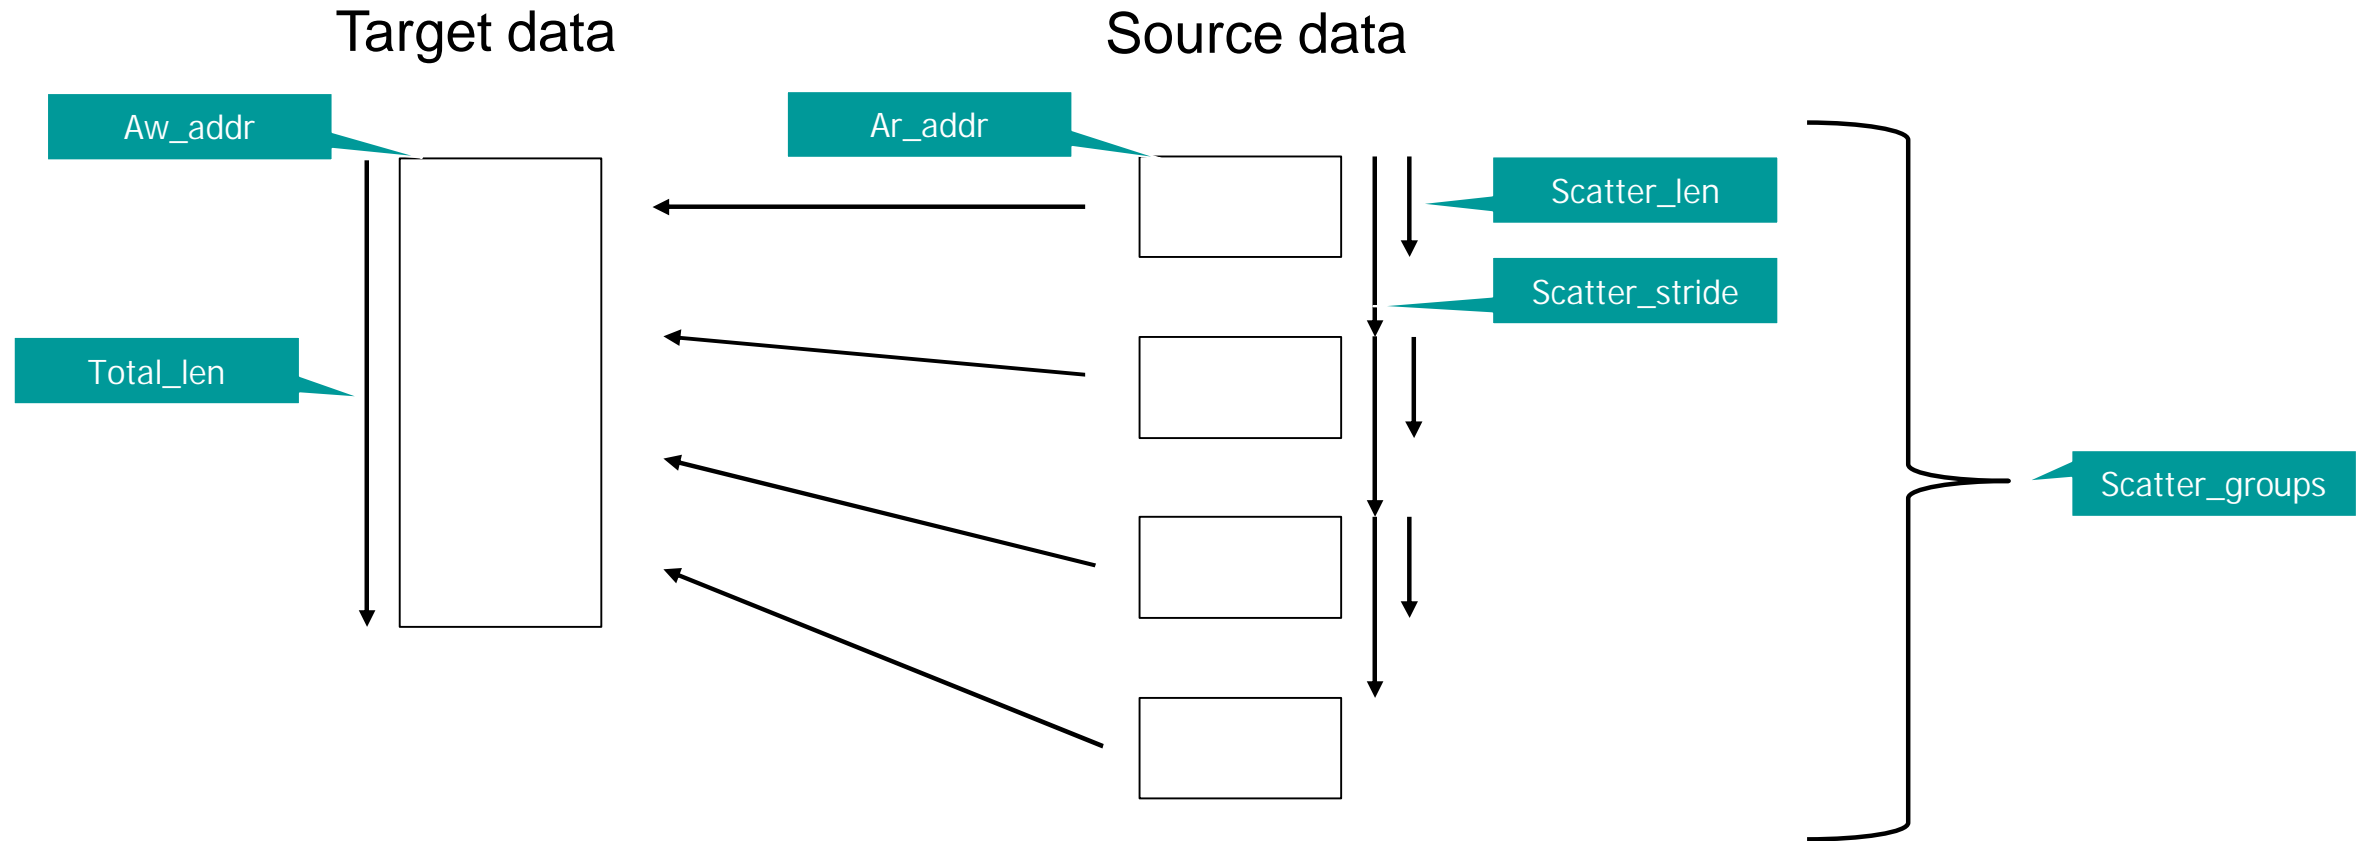


All addresses are byte addresses

All lengths are byte lengths

$\text{total_len} = \text{scatter_groups} * \text{scatter_len}$

Gather Operation

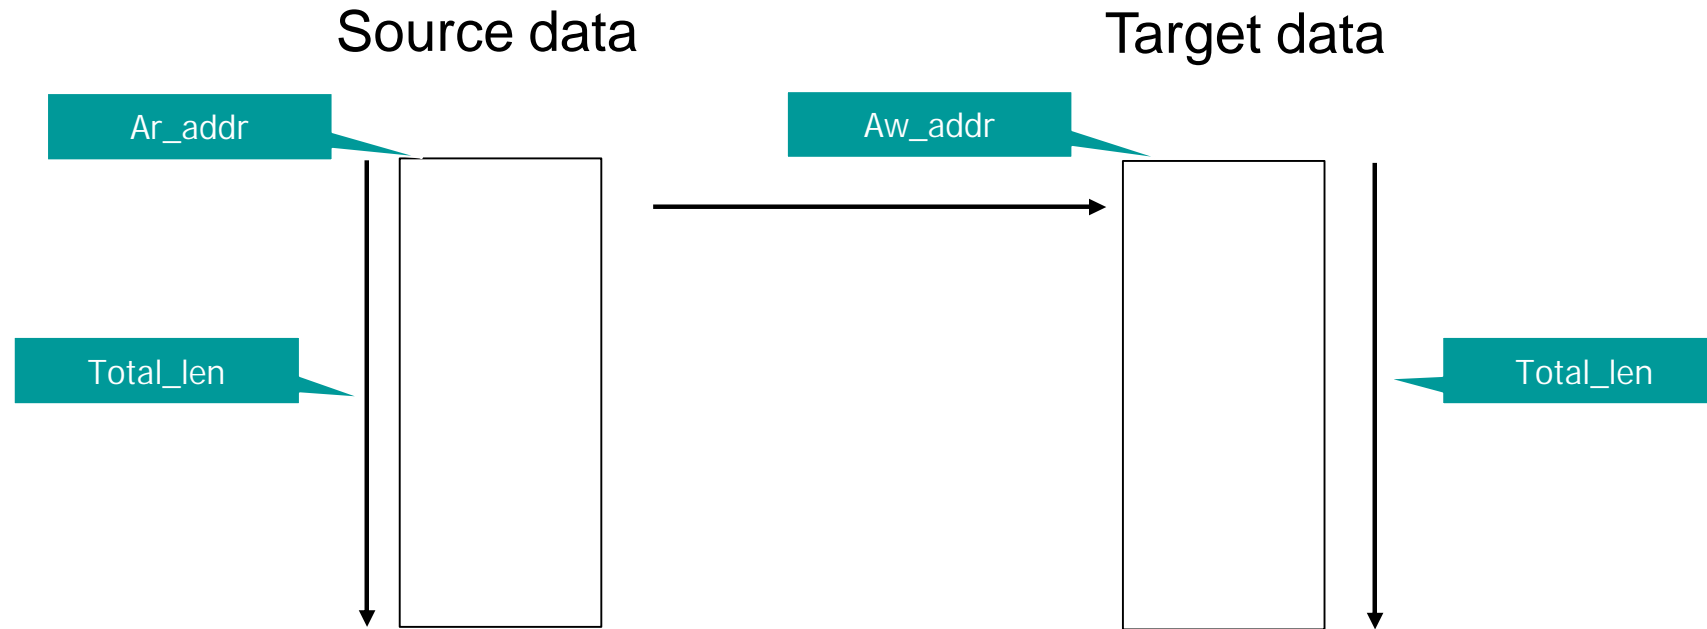


All addresses are byte addresses

All lengths are byte lengths

$\text{total_len} = \text{scatter_groups} * \text{scatter_len}$

Copy Operation



All addresses are byte addresses

All lengths are byte lengths

For copy operation, scatter_len, scatter_stride, scatter_groups are ignored

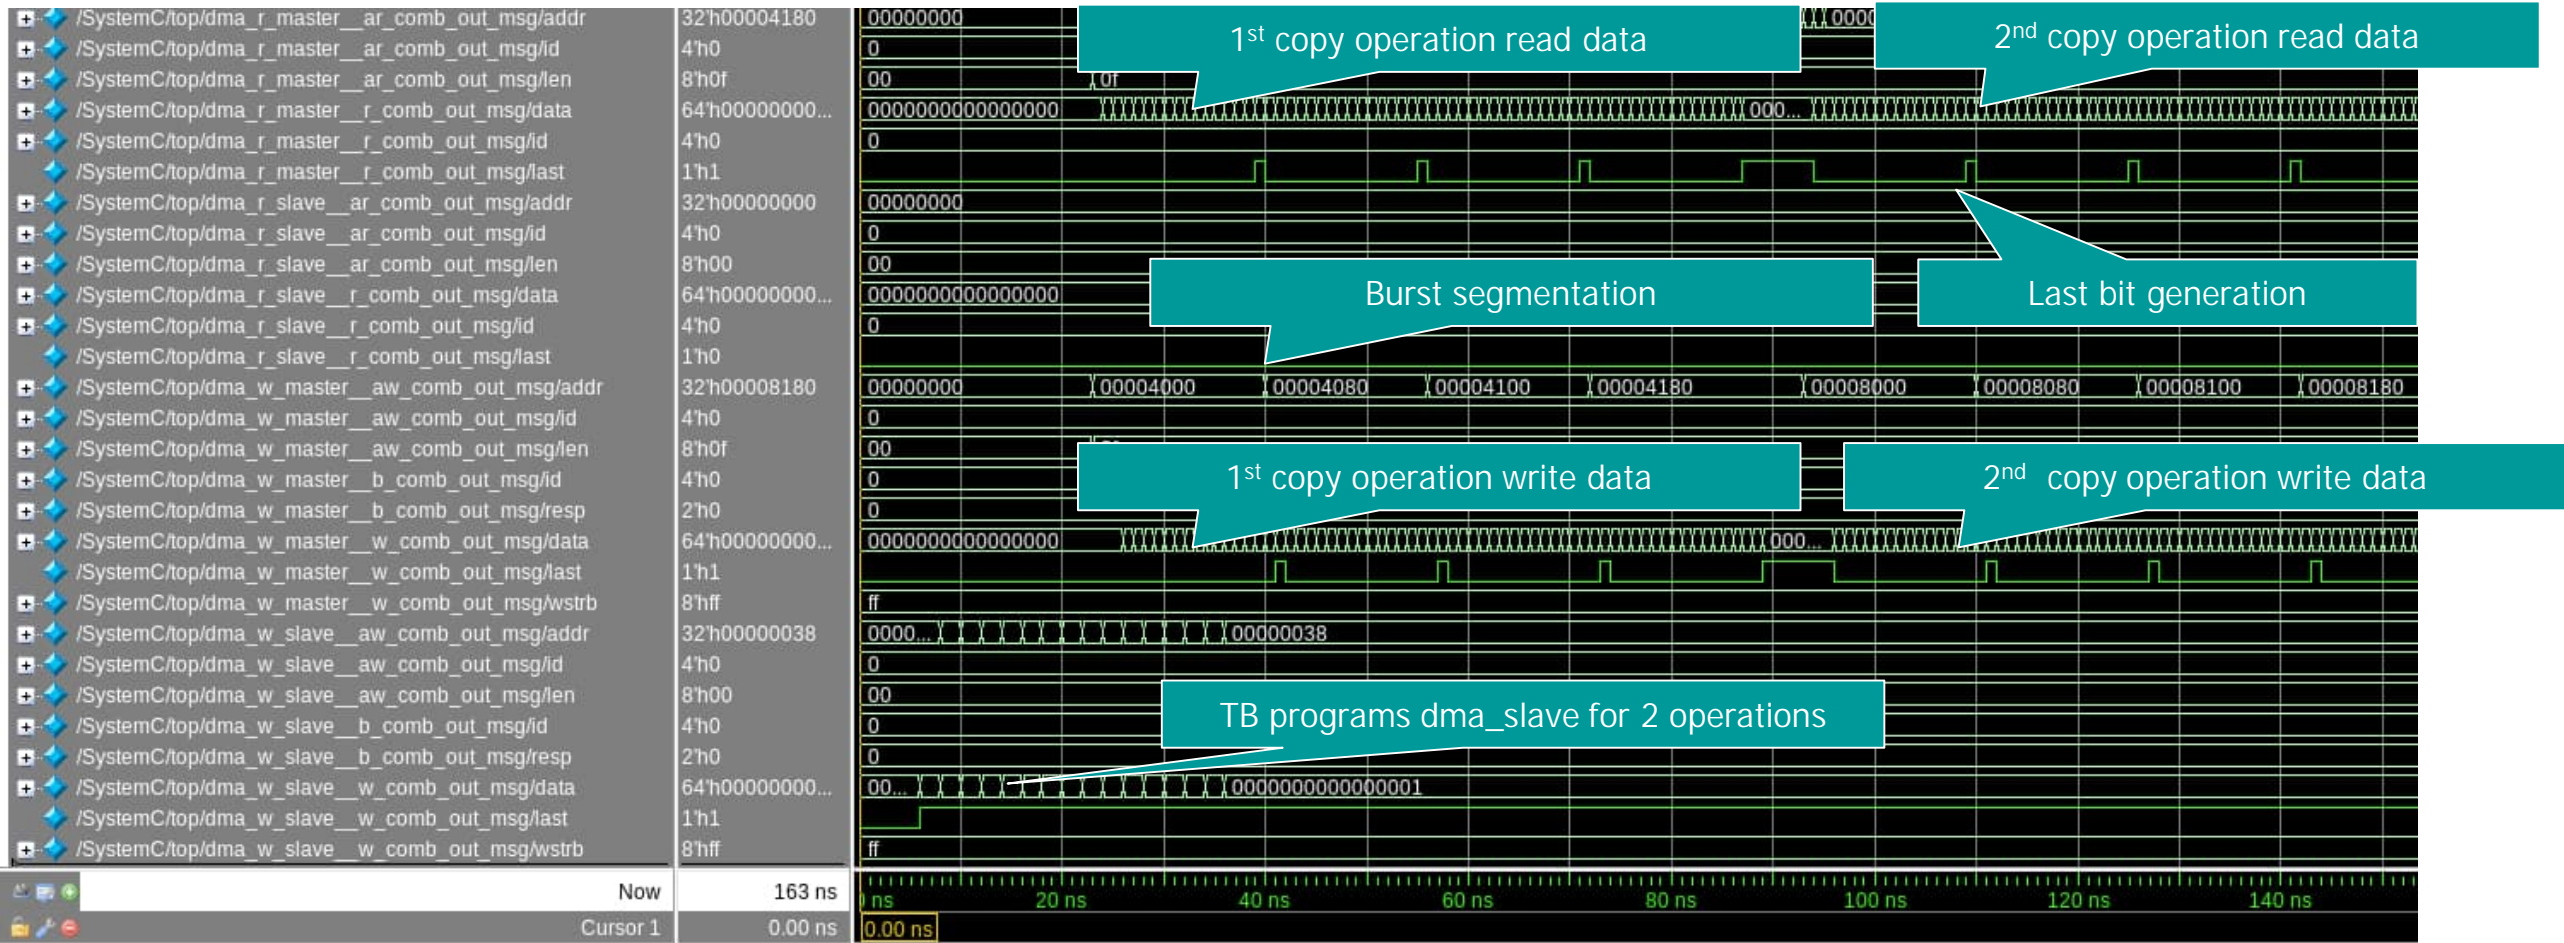
SC Testbench

```
62 int test_iterations = 1;  
63 bool copy_mode = true;  
64 int total_len = 64 * bytesPerBeat;  
65 int scatter_groups = 4;  
66 int scatter_len = total_len / scatter_groups;  
67 int scatter_stride = scatter_len * 2;  
68 int source_addr = 0x1000;  
69 int target1_addr = 0x4000;  
70 int target2_addr = 0x8000;  
71 sc_time start_time, end_time;  
72
```

If true, each iteration is a copy then a copy, else it is a scatter then a gather


Target address for first operation, source address for second operation

Copy_mode = true, test_iterations = 1

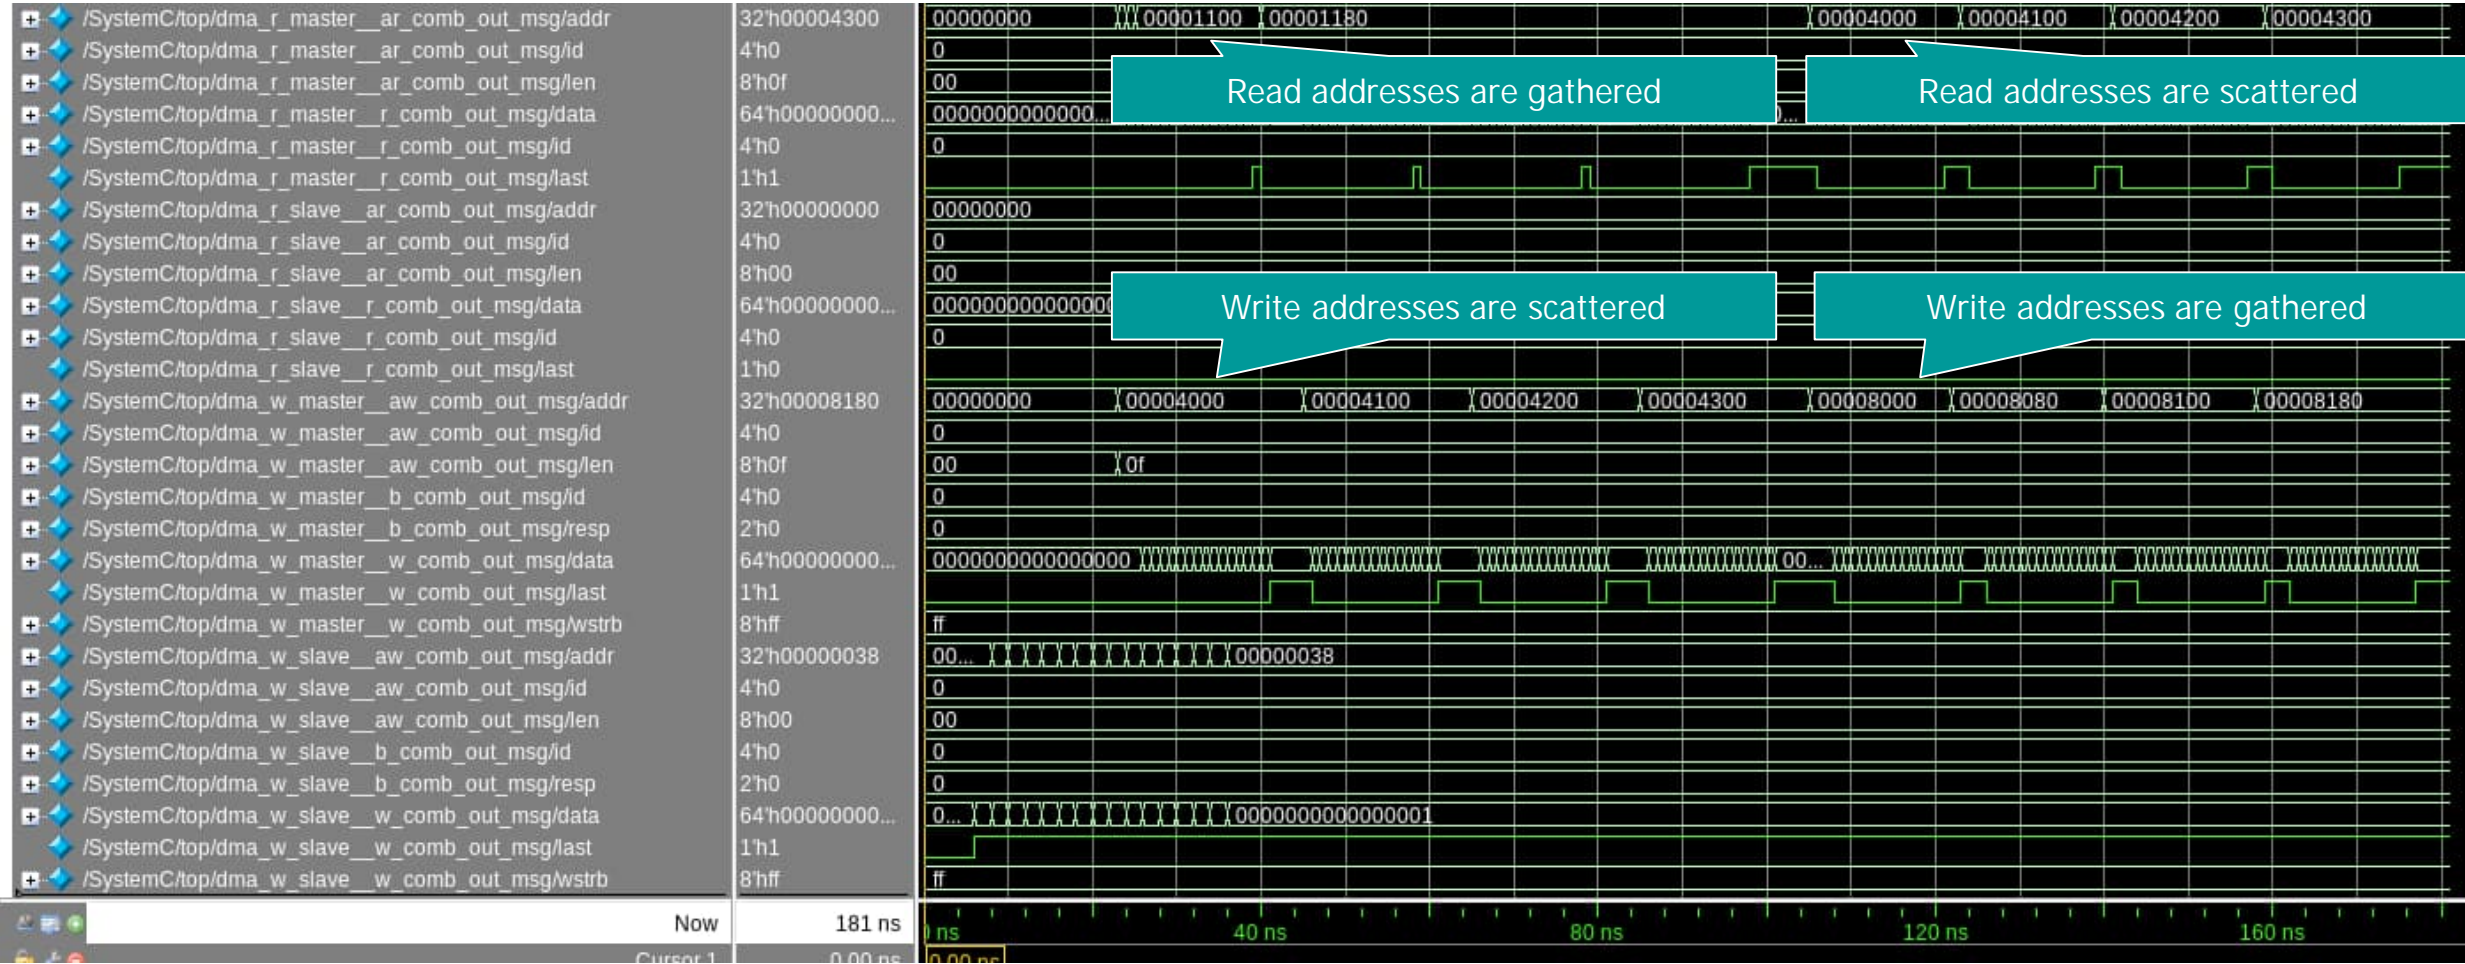


Note: SystemC Waveforms, (not RTL) Makefile sets segmentation size to 16 rather than 256 so easier to see in waveforms

Same scenario, but in RTL

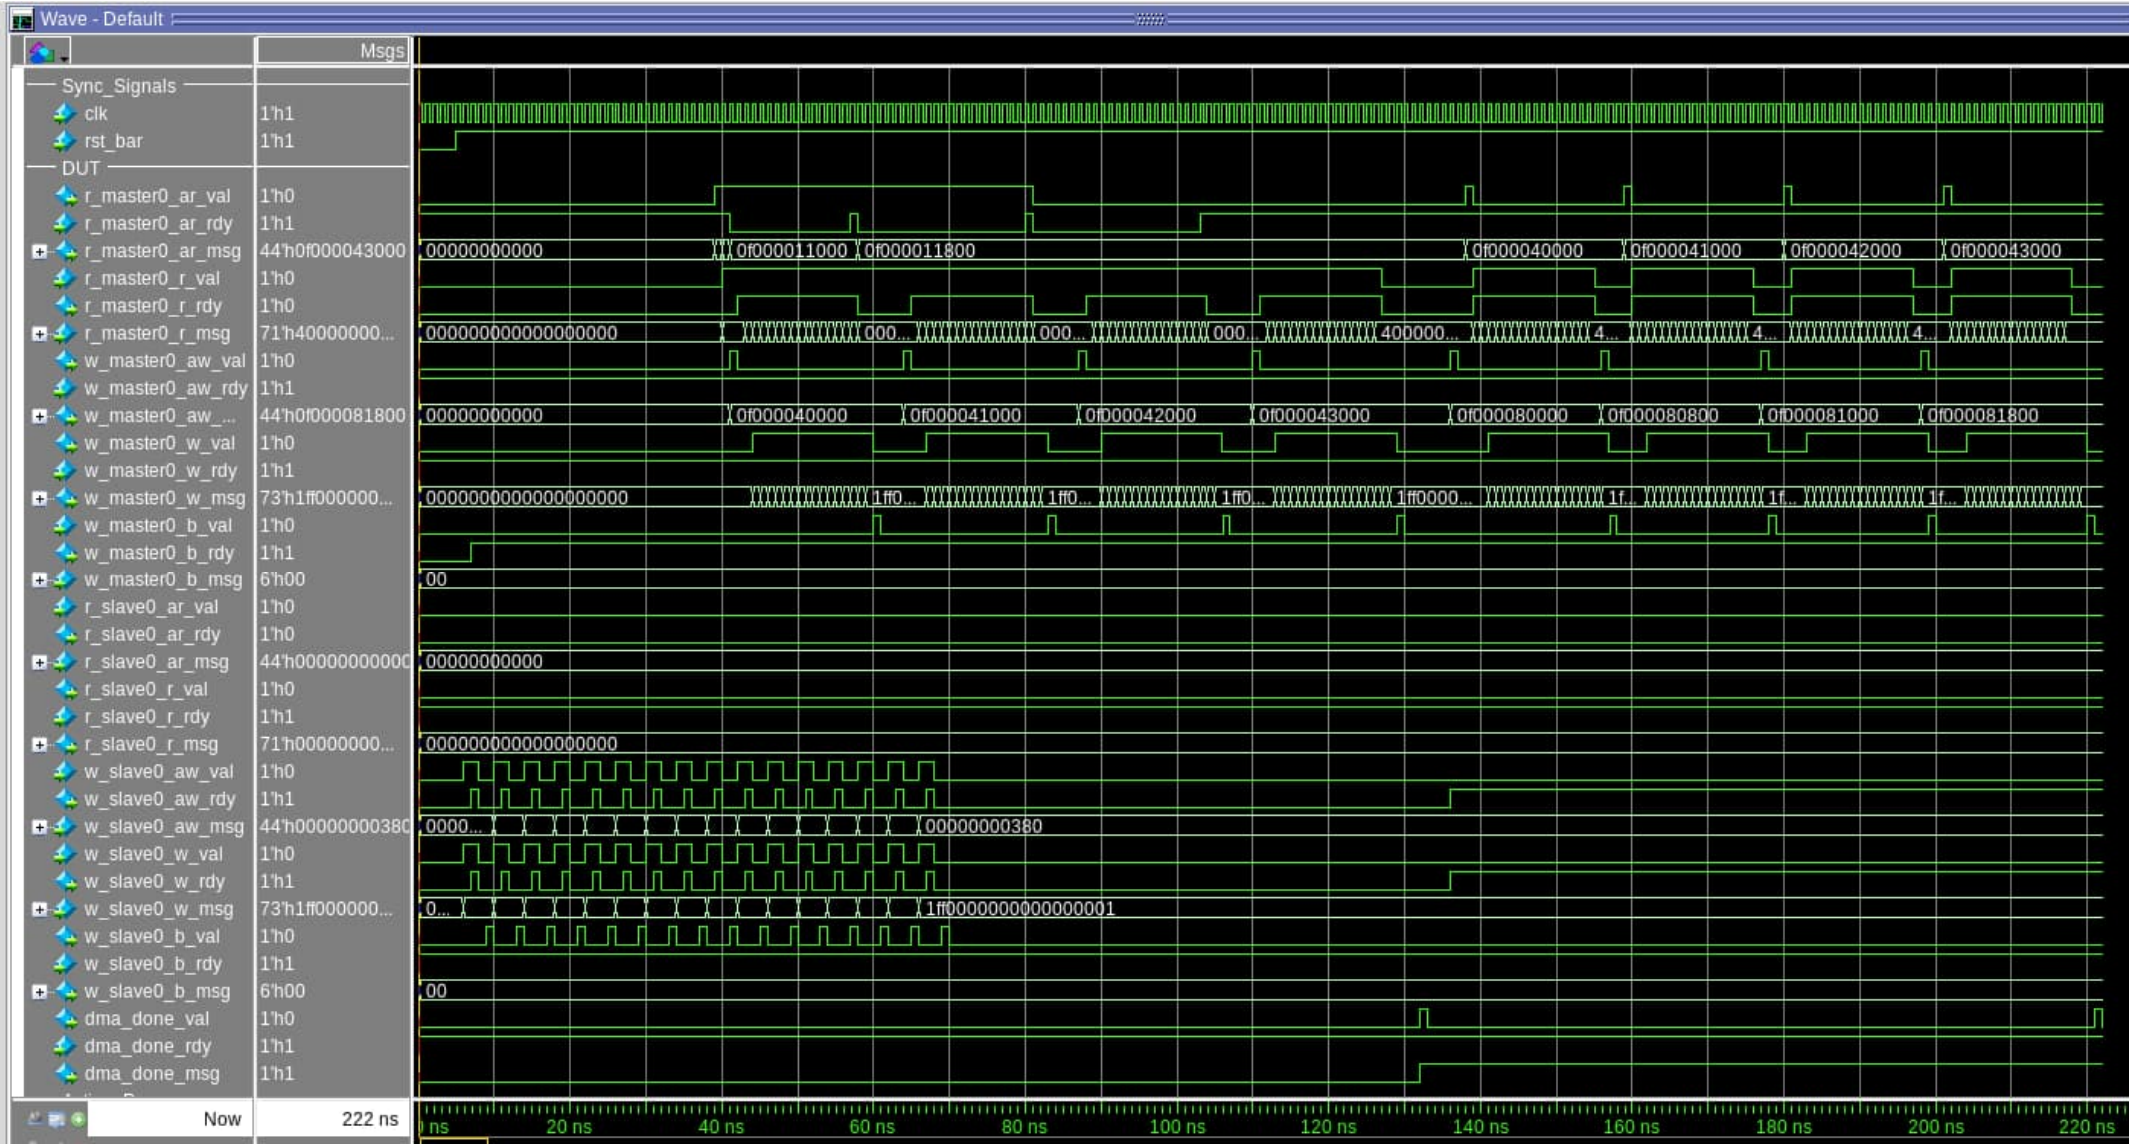


Copy_mode = false, test_iterations = 1



Note: SystemC Waveforms, (not RTL)

Same scenario, but in RTL



Keeping things simple..

To simplify both the DUT and the TB, the DUT enforces that all addresses and lengths are aligned to bus datawidth boundaries

When TB writes to “start” CSR, checks are enforced and a AXI4 error response is returned if there any errors.

```
295 case offsetof(dma_address_map, start):
296     if ((cmd1.ar_addr & (bytesPerBeat-1)) ||
297         (cmd1.aw_addr & (bytesPerBeat-1)) ||
298         (cmd1.total_len & (bytesPerBeat-1)))
299     {
300         LOG("discarding invalid DMA command");
301         break;
302     }
303     if (cmd1.dma_mode != dma_mode_t::COPY)
304         if ((cmd1.scatter_len & (bytesPerBeat-1)) ||
305             (cmd1.scatter_stride & (bytesPerBeat-1)) ||
306             (cmd1.scatter_len * cmd1.scatter_groups != cmd1.total_len))
307     {
308         LOG("discarding invalid DMA command");
309         break;
310     }
311     dma_cmd_chan.Push(cmd1);
312     b.resp = Enc::XRESP::OKAY;
313     break;
314
315 default:
316     break;
317 }
318
319 w_slave0.b.Push(b);
```

When TB writes to “start” CSR

Make sure addresses and lengths are aligned

Do additional checks if not in COPY mode

Send AXI4 error code to TB if any checks fail

COPY implementation in DMA

```

122 while(1) {
123     dma_cmd cmd = dma_cmd_chan.Pop();
124
125     switch (cmd.dma_mode) {
126     case dma_mode_t::COPY: {
127         ex_ar_payload ar;
128         ex_aw_payload aw;
129         ar.ex_len = (cmd.total_len / bytesPerBeat) - 1;
130         aw.ex_len = (cmd.total_len / bytesPerBeat) - 1;
131         ar.addr = cmd.ar_addr;
132         aw.addr = cmd.aw_addr;
133         r_segment0_ex_ar_chan.Push(ar);
134         w_segment0_ex_aw_chan.Push(aw);
135
136         #pragma hls_pipeline_init_interval 1
137         #pragma pipeline_stall_mode flush
138         while (1) {
139             r_payload r = r_master0.r.Pop();
140             w_payload w;
141             w.data = r.data;
142             w_segment0_w_chan.Push(w);
143
144             if (ar.ex_len-- == 0)
145                 break;
146         }
147
148         b_payload b;
149         b = w_segment0_b_chan.Pop();
150         dma_done.Push(true);
151         break;
152     }
153 }

```

Pop next DMA command


Convert byte length to beat length (AXI4 beat length encoding)

Push out AR and AW bursts

Copy data from R to W

We are done when length is currently 0 (AXI4 beat encoding)

Pop write response from b channel



SCATTER implementation in DMA

```
154 case dma_mode_t::SCATTER: {
155     ex_ar_payload ar;
156     ex_aw_payload aw;
157     ar.ex_len = (cmd.total_len / bytesPerBeat) - 1;
158     aw.ex_len = (cmd.scatter_len / bytesPerBeat) - 1;
159     ar.addr = cmd.ar_addr;
160     aw.addr = cmd.aw_addr;
161     r_segment0_ex_ar_chan.Push(ar);
162
163     while (1) {
164         w_segment0_ex_aw_chan.Push(aw);
165
166         #pragma hls_pipeline_init_interval 1
167         #pragma pipeline_stall_mode flush
168         while (1) {
169             r_payload r = r_master0.r.Pop();
170             w_payload w;
171             w.data = r.data;
172             w_segment0_w_chan.Push(w);
173
174             if (aw.ex_len-- == 0)
175                 break;
176         }
177
178         w_segment0_b_chan.Pop();
179         aw.addr += cmd.scatter_stride;
180         aw.ex_len = (cmd.scatter_len / bytesPerBeat) - 1;
181
182         cmd.total_len -= cmd.scatter_len;
183         if (cmd.total_len == 0)
184             break;
185     }
186
187     dma_done.Push(true);
188     break;
189 }
```

Convert byte length to beat length (AXI4 beat length encoding)

aw.ex_len is computed based on scatter_len, not total_len

Start the AR burst, but not the AW yet

Start the AW burst for current region in the scatter groups


Copy data from R to W

We are done when length is currently 0 (AXI4 beat encoding)

Need to do b.Pop for every Push(aw)

Increment aw.addr for next scatter group region

We are done when the remaining total_len is zero



Homework #1 - GATHER implementation in DMA

Steps:

1. Edit testbench.cpp and change copy_mode to be false
2. Edit scatter_gather.h and write the code for GATHER
 - This is strictly confined to the GATHER branch of the case statement.
3. Compile and run your SC code and make sure the TB self-check passes
4. View SC and RTL waveforms – see README file in same dir.

Homework #2 – Optimize the beat rate

- For `copy_mode = false` and `test_iterations = 10`, make small modifications to model and synthesis directives to optimize throughput (as reported in log output as "beat rate").
- Measure in both SC sim and RTL sim.